

```
In [1]: import pandas as pd

In [2]: applicationdf = pd.read_csv(r"C:\Users\godis\Desktop\OneDrive - University of St

In [3]: applicationdf.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 307511 entries, 0 to 307510
Columns: 122 entries, SK_ID_CURR to AMT_REQ_CREDIT_BUREAU_YEAR
dtypes: float64(65), int64(41), object(16)
memory usage: 286.2+ MB

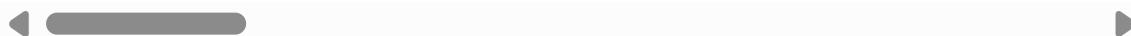
In [4]: print(applicationdf.isnull().sum())

SK_ID_CURR          0
TARGET              0
NAME_CONTRACT_TYPE 0
CODE_GENDER         0
FLAG_OWN_CAR        0
...
AMT_REQ_CREDIT_BUREAU_DAY 41519
AMT_REQ_CREDIT_BUREAU_WEEK 41519
AMT_REQ_CREDIT_BUREAU_MON 41519
AMT_REQ_CREDIT_BUREAU_QRT 41519
AMT_REQ_CREDIT_BUREAU_YEAR 41519
Length: 122, dtype: int64
```

```
In [5]: applicationdf.head()
```

	SK_ID_CURR	TARGET	NAME_CONTRACT_TYPE	CODE_GENDER	FLAG_OWN_CAR	FLA
0	100002	1	Cash loans	M	N	
1	100003	0	Cash loans	F	N	
2	100004	0	Revolving loans	M	Y	
3	100006	0	Cash loans	F	N	
4	100007	0	Cash loans	M	N	

5 rows × 122 columns



```
In [6]: # data loaded in data frame
missing_percentage = applicationdf.isnull().sum() * 100 / len(applicationdf)

# Identifying columns where more than 50% of the data is missing
high_missing_cols = missing_percentage[missing_percentage > 50].sort_values(ascending=True)

# Storing the list of columns to drop
cols_to_drop = high_missing_cols.index.tolist()

# Dropping the columns
applicationdf.drop(columns=cols_to_drop, inplace=True)

# checking whether columns dropped or not
```

```
print("Columns Dropped (over 50% missing data):", cols_to_drop)
print(f"New DataFrame shape: {applicationdf.shape}")
```

```
Columns Dropped (over 50% missing data): ['COMMONAREA_AVG', 'COMMONAREA_MEDI', 'COMMONAREA_MODE', 'NONLIVINGAPARTMENTS_AVG', 'NONLIVINGAPARTMENTS_MODE', 'NONLIVINGAPARTMENTS_MEDI', 'FONDKAPREMONT_MODE', 'LIVINGAPARTMENTS_AVG', 'LIVINGAPARTMENTS_MODE', 'LIVINGAPARTMENTS_MEDI', 'FLOORSMIN_MODE', 'FLOORSMIN_MEDI', 'FLOORSMIN_AVG', 'YEARS_BUILD_MODE', 'YEARS_BUILD_MEDI', 'YEARS_BUILD_AVG', 'OWN_CAR_AGE', 'LANDAREA_MODE', 'LANDAREA_MEDI', 'LANDAREA_AVG', 'BASEMENTAREA_MODE', 'BASEMENTAREA_AVG', 'BASEMENTAREA_MEDI', 'EXT_SOURCE_1', 'NONLIVINGAREA_MEDI', 'NONLIVINGAREA_AVG', 'NONLIVINGAREA_MODE', 'ELEVATORS_AVG', 'ELEVATORS_MODE', 'ELEVATORS_MEDI', 'WALLSMATERIAL_MODE', 'APARTMENTS_MEDI', 'APARTMENTS_AVG', 'APARTMENTS_MODE', 'ENTRANCES_AVG', 'ENTRANCES_MODE', 'ENTRANCES_MEDI', 'LIVINGAREA_MODE', 'LIVINGAREA_AVG', 'LIVINGAREA_MEDI', 'HOUSETYPE_MODE']  
New DataFrame shape: (307511, 81)
```

In [7]: `applicationdf.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 307511 entries, 0 to 307510
Data columns (total 81 columns):
 #   Column           Non-Null Count Dtype
 --- 
 0   SK_ID_CURR       307511 non-null int64
 1   TARGET           307511 non-null int64
 2   NAME_CONTRACT_TYPE 307511 non-null object
 3   CODE_GENDER       307511 non-null object
 4   FLAG_OWN_CAR      307511 non-null object
 5   FLAG_OWN_REALTY   307511 non-null object
 6   CNT_CHILDREN      307511 non-null int64
 7   AMT_INCOME_TOTAL  307511 non-null float64
 8   AMT_CREDIT         307511 non-null float64
 9   AMT_ANNUITY        307499 non-null float64
 10  AMT_GOODS_PRICE   307233 non-null float64
 11  NAME_TYPE_SUITE   306219 non-null object
 12  NAME_INCOME_TYPE  307511 non-null object
 13  NAME_EDUCATION_TYPE 307511 non-null object
 14  NAME_FAMILY_STATUS 307511 non-null object
 15  NAME_HOUSING_TYPE 307511 non-null object
 16  REGION_POPULATION_RELATIVE 307511 non-null float64
 17  DAYS_BIRTH         307511 non-null int64
 18  DAYS_EMPLOYED      307511 non-null int64
 19  DAYS_REGISTRATION  307511 non-null float64
 20  DAYS_ID_PUBLISH   307511 non-null int64
 21  FLAG_MOBIL         307511 non-null int64
 22  FLAG_EMP_PHONE     307511 non-null int64
 23  FLAG_WORK_PHONE    307511 non-null int64
 24  FLAG_CONT_MOBILE   307511 non-null int64
 25  FLAG_PHONE          307511 non-null int64
 26  FLAG_EMAIL          307511 non-null int64
 27  OCCUPATION_TYPE    211120 non-null object
 28  CNT_FAM_MEMBERS    307509 non-null float64
 29  REGION_RATING_CLIENT 307511 non-null int64
 30  REGION_RATING_CLIENT_W_CITY 307511 non-null int64
 31  WEEKDAY_APPR_PROCESS_START 307511 non-null object
 32  HOUR_APPR_PROCESS_START 307511 non-null int64
 33  REG_REGION_NOT_LIVE_REGION 307511 non-null int64
 34  REG_REGION_NOT_WORK_REGION 307511 non-null int64
 35  LIVE_REGION_NOT_WORK_REGION 307511 non-null int64
 36  REG_CITY_NOT_LIVE_CITY 307511 non-null int64
 37  REG_CITY_NOT_WORK_CITY 307511 non-null int64
 38  LIVE_CITY_NOT_WORK_CITY 307511 non-null int64
 39  ORGANIZATION_TYPE   307511 non-null object
 40  EXT_SOURCE_2        306851 non-null float64
 41  EXT_SOURCE_3        246546 non-null float64
 42  YEARS_BEGINEXPLUATATION_AVG 157504 non-null float64
 43  FLOORSMAX_AVG       154491 non-null float64
 44  YEARS_BEGINEXPLUATATION_MODE 157504 non-null float64
 45  FLOORSMAX_MODE      154491 non-null float64
 46  YEARS_BEGINEXPLUATATION_MEDI 157504 non-null float64
 47  FLOORSMAX_MEDI      154491 non-null float64
 48  TOTALAREA_MODE       159080 non-null float64
 49  EMERGENCYSTATE_MODE 161756 non-null object
 50  OBS_30_CNT_SOCIAL_CIRCLE 306490 non-null float64
 51  DEF_30_CNT_SOCIAL_CIRCLE 306490 non-null float64
 52  OBS_60_CNT_SOCIAL_CIRCLE 306490 non-null float64
 53  DEF_60_CNT_SOCIAL_CIRCLE 306490 non-null float64
 54  DAYS_LAST_PHONE_CHANGE 307510 non-null float64
```

```

55 FLAG_DOCUMENT_2           307511 non-null int64
56 FLAG_DOCUMENT_3           307511 non-null int64
57 FLAG_DOCUMENT_4           307511 non-null int64
58 FLAG_DOCUMENT_5           307511 non-null int64
59 FLAG_DOCUMENT_6           307511 non-null int64
60 FLAG_DOCUMENT_7           307511 non-null int64
61 FLAG_DOCUMENT_8           307511 non-null int64
62 FLAG_DOCUMENT_9           307511 non-null int64
63 FLAG_DOCUMENT_10          307511 non-null int64
64 FLAG_DOCUMENT_11          307511 non-null int64
65 FLAG_DOCUMENT_12          307511 non-null int64
66 FLAG_DOCUMENT_13          307511 non-null int64
67 FLAG_DOCUMENT_14          307511 non-null int64
68 FLAG_DOCUMENT_15          307511 non-null int64
69 FLAG_DOCUMENT_16          307511 non-null int64
70 FLAG_DOCUMENT_17          307511 non-null int64
71 FLAG_DOCUMENT_18          307511 non-null int64
72 FLAG_DOCUMENT_19          307511 non-null int64
73 FLAG_DOCUMENT_20          307511 non-null int64
74 FLAG_DOCUMENT_21          307511 non-null int64
75 AMT_REQ_CREDIT_BUREAU_HOUR 265992 non-null float64
76 AMT_REQ_CREDIT_BUREAU_DAY 265992 non-null float64
77 AMT_REQ_CREDIT_BUREAU_WEEK 265992 non-null float64
78 AMT_REQ_CREDIT_BUREAU_MON 265992 non-null float64
79 AMT_REQ_CREDIT_BUREAU_QRT 265992 non-null float64
80 AMT_REQ_CREDIT_BUREAU_YEAR 265992 non-null float64
dtypes: float64(27), int64(41), object(13)
memory usage: 190.0+ MB

```

```

In [8]: # Finding all remaining columns that are numerical (float or int) AND have missing values
# Re-running the original check on the updated DataFrame
numerical_missing_cols = applicationdf.select_dtypes(include=['float64', 'int64'])

print("--- Numerical Columns to Impute ---")
print(numerical_missing_cols)
print("-----")

for col in numerical_missing_cols:
    median_val = applicationdf[col].median()
    applicationdf[col].fillna(median_val, inplace=True)
    print(f"Imputed numerical column '{col}' with median: {median_val}")

# Verifying all numerical missing values are resolved or not
print("\nVerification: Remaining missing values in numerical columns:")
print(applicationdf[numerical_missing_cols].isnull().sum().sum())

```

```

--- Numerical Columns to Impute ---
['AMT_ANNUITY', 'AMT_GOODS_PRICE', 'CNT_FAM_MEMBERS', 'EXT_SOURCE_2', 'EXT_SOURCE_3', 'YEARS_BEGINEXPLUATATION_AVG', 'FLOORSMAX_AVG', 'YEARS_BEGINEXPLUATATION_MODE', 'FLOORSMAX_MODE', 'YEARS_BEGINEXPLUATATION_MEDI', 'FLOORSMAX_MEDI', 'TOTALAREAMODE', 'OBS_30_CNT_SOCIAL_CIRCLE', 'DEF_30_CNT_SOCIAL_CIRCLE', 'OBS_60_CNT_SOCIAL_CIRCLE', 'DEF_60_CNT_SOCIAL_CIRCLE', 'DAYS_LAST_PHONE_CHANGE', 'AMT_REQ_CREDIT_BUREAU_HOUR', 'AMT_REQ_CREDIT_BUREAU_DAY', 'AMT_REQ_CREDIT_BUREAU_WEEK', 'AMT_REQ_CREDIT_BUREAU_MON', 'AMT_REQ_CREDIT_BUREAU_QRT', 'AMT_REQ_CREDIT_BUREAU_YEAR']

-----
Imputed numerical column 'AMT_ANNUITY' with median: 24903.0
Imputed numerical column 'AMT_GOODS_PRICE' with median: 450000.0
Imputed numerical column 'CNT_FAM_MEMBERS' with median: 2.0
Imputed numerical column 'EXT_SOURCE_2' with median: 0.5659614260608526
Imputed numerical column 'EXT_SOURCE_3' with median: 0.5352762504724826
Imputed numerical column 'YEARS_BEGINEXPLUATATION_AVG' with median: 0.9816
Imputed numerical column 'FLOORSMAX_AVG' with median: 0.1667
Imputed numerical column 'YEARS_BEGINEXPLUATATION_MODE' with median: 0.9816
Imputed numerical column 'FLOORSMAX_MODE' with median: 0.1667
Imputed numerical column 'YEARS_BEGINEXPLUATATION_MEDI' with median: 0.9816
Imputed numerical column 'FLOORSMAX_MEDI' with median: 0.1667
Imputed numerical column 'TOTALAREA_MODE' with median: 0.0688
Imputed numerical column 'OBS_30_CNT_SOCIAL_CIRCLE' with median: 0.0
Imputed numerical column 'DEF_30_CNT_SOCIAL_CIRCLE' with median: 0.0
Imputed numerical column 'OBS_60_CNT_SOCIAL_CIRCLE' with median: 0.0
Imputed numerical column 'DEF_60_CNT_SOCIAL_CIRCLE' with median: 0.0
Imputed numerical column 'DAYS_LAST_PHONE_CHANGE' with median: -757.0
Imputed numerical column 'AMT_REQ_CREDIT_BUREAU_HOUR' with median: 0.0
Imputed numerical column 'AMT_REQ_CREDIT_BUREAU_DAY' with median: 0.0
Imputed numerical column 'AMT_REQ_CREDIT_BUREAU_WEEK' with median: 0.0
Imputed numerical column 'AMT_REQ_CREDIT_BUREAU_MON' with median: 0.0
Imputed numerical column 'AMT_REQ_CREDIT_BUREAU_QRT' with median: 0.0
Imputed numerical column 'AMT_REQ_CREDIT_BUREAU_YEAR' with median: 1.0

```

Verification: Remaining missing values in numerical columns:

0

C:\Users\godis\AppData\Local\Temp\ipykernel_13044\1207936534.py:11: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an `inplace` method.

The behavior will change in pandas 3.0. This `inplace` method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing `'df[col].method(value, inplace=True)'`, try using `'df.method({col: value}, inplace=True)'` or `df[col] = df[col].method(value)` instead, to perform the operation `inplace` on the original object.

```
applicationdf[col].fillna(median_val, inplace=True)
```

```
In [9]: # Identifying all remaining categorical columns with missing values
categorical_missing_cols = applicationdf.select_dtypes(include=['object']).columns

print("\n--- Categorical Columns to Impute ---")
print(categorical_missing_cols)
print("-----")

for col in categorical_missing_cols:
    # Filling missing values with 'Missing'
    applicationdf[col].fillna('Missing', inplace=True)
    print(f"Imputed categorical column '{col}' with 'Missing'")
```

```
# Verifying that all missing values in the entire DataFrame are resolved
print(f"\nTotal remaining missing values in DataFrame: {applicationdf.isnull().sum().sum()}\n")

--- Categorical Columns to Impute ---
['NAME_TYPE_SUITE', 'OCCUPATION_TYPE', 'EMERGENCYSTATE_MODE']

-----
Imputed categorical column 'NAME_TYPE_SUITE' with 'Missing'
Imputed categorical column 'OCCUPATION_TYPE' with 'Missing'
Imputed categorical column 'EMERGENCYSTATE_MODE' with 'Missing'
```

Total remaining missing values in DataFrame: 0

C:\Users\godis\AppData\Local\Temp\ipykernel_13044\1702331462.py:10: FutureWarning:
A value is trying to be set on a copy of a DataFrame or Series through chained
assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
applicationdf[col].fillna('Missing', inplace=True)
```

```
In [10]: # Identifying all the object columns that need encoding
categorical_cols_to_encode = applicationdf.select_dtypes(include=['object']).columns

# Applying One-Hot Encoding and dropping the original categorical columns
applicationdf = pd.get_dummies(applicationdf, columns=categorical_cols_to_encode)

# Verifying the final shape
print(f"\nFinal DataFrame shape after encoding: {applicationdf.shape}")
print(f"Number of columns after encoding: {applicationdf.shape[1]}")

print("Sample Encoded Columns (First 5 Rows):")
print(applicationdf.filter(regex='^M$|^F$|Cash_loans$|Missing$').head())
```

Final DataFrame shape after encoding: (307511, 184)

Number of columns after encoding: 184

Sample Encoded Columns (First 5 Rows):

	CODE_GENDER_M	NAME_TYPE_SUITE_Missing	OCCUPATION_TYPE_Missing
0	True	False	False
1	False	False	False
2	True	False	False
3	False	False	False
4	True	False	False

```
In [11]: import matplotlib.pyplot as plt
import seaborn as sns

# Setting a style for better visualization
plt.style.use("seaborn-v0_8-darkgrid")

# 1. Checking the distribution of the TARGET variable
target_counts = applicationdf['TARGET'].value_counts()
target_proportions = target_counts / len(applicationdf) * 100
```

```

print("Target Variable Distribution:")
print(target_counts)
print(f"Proportion of Non-Defaulters (0): {target_proportions.iloc[0]:.2f}%")
print(f"Proportion of Defaulters (1): {target_proportions.iloc[1]:.2f}%")

# 2. Visualizing the distribution using a Bar Chart
plt.figure(figsize=(7, 5))
sns.barplot(x=target_counts.index, y=target_counts.values, palette=['skyblue', 'salmon'])
plt.title('Distribution of Loan Status (TARGET)', fontsize=14)
plt.xlabel('Loan Status (0: Repaid, 1: Defaulted)')
plt.ylabel('Number of Applicants')
plt.xticks([0, 1])
plt.show()

```

Target Variable Distribution:

TARGET

0 282686

1 24825

Name: count, dtype: int64

Proportion of Non-Defaulters (0): 91.93%

Proportion of Defaulters (1): 8.07%

C:\Users\godis\AppData\Local\Temp\ipykernel_13044\4138747666.py:18: FutureWarning:

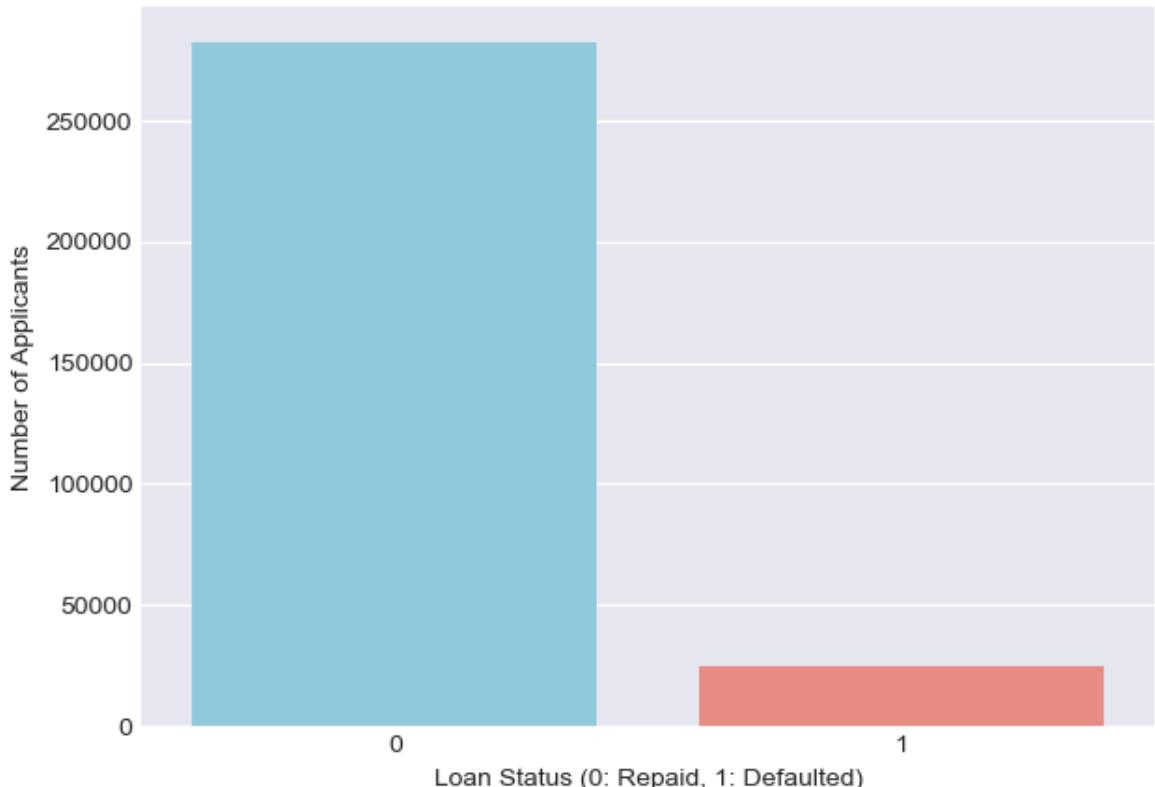
Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```

sns.barplot(x=target_counts.index, y=target_counts.values, palette=['skyblue', 'salmon'])

```

Distribution of Loan Status (TARGET)



In [12]: # Selecting the most critical financial and demographic features
key_numeric_features = ['AMT_INCOME_TOTAL', 'AMT_CREDIT', 'AMT_ANNUITY', 'DAYS_B

```

# 1. Descriptive Statistics Table
print("\n--- Descriptive Statistics Table ---")
print(applicationdf[key_numeric_features].describe())

# 2. Box Plots
plt.figure(figsize=(15, 6))
applicationdf[key_numeric_features].boxplot()
plt.title('Box Plots for Key Numerical Features (Outlier Check)', fontsize=14)
plt.ylabel('Value')
plt.show()

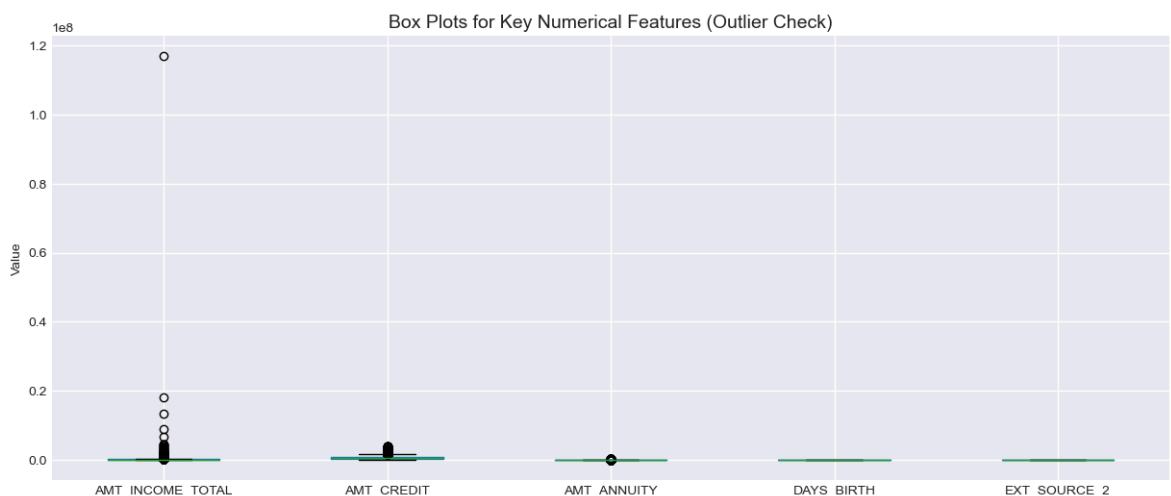
# 3. Histograms
plt.figure(figsize=(15, 4))
for i, col in enumerate(key_numeric_features[:3]): # Focusing on the monetary columns
    plt.subplot(1, 3, i + 1)
    sns.histplot(applicationdf[col], kde=True, bins=50, color='teal')
    plt.title(f'Distribution of {col}', fontsize=12)
    plt.xlabel(col)
plt.tight_layout()
plt.show()

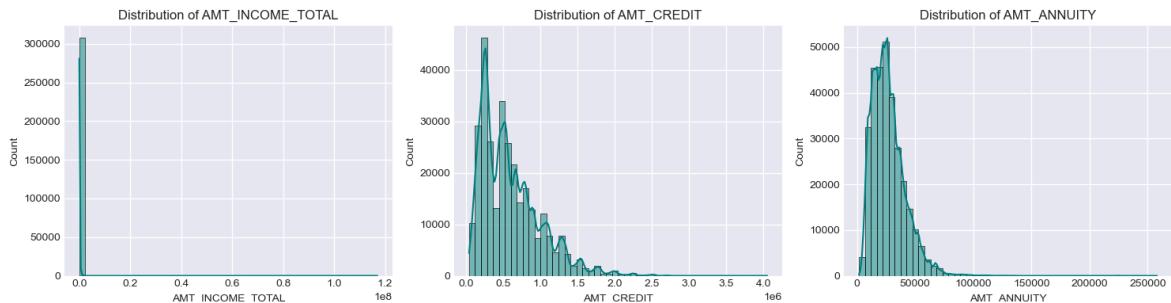
```

--- Descriptive Statistics Table ---

	AMT_INCOME_TOTAL	AMT_CREDIT	AMT_ANNUITY	DAYS_BIRTH	\
count	3.075110e+05	3.075110e+05	307511.000000	307511.000000	
mean	1.687979e+05	5.990260e+05	27108.487841	-16036.995067	
std	2.371231e+05	4.024908e+05	14493.461065	4363.988632	
min	2.565000e+04	4.500000e+04	1615.500000	-25229.000000	
25%	1.125000e+05	2.700000e+05	16524.000000	-19682.000000	
50%	1.471500e+05	5.135310e+05	24903.000000	-15750.000000	
75%	2.025000e+05	8.086500e+05	34596.000000	-12413.000000	
max	1.170000e+08	4.050000e+06	258025.500000	-7489.000000	

	EXT_SOURCE_2
count	3.075110e+05
mean	5.145034e-01
std	1.908699e-01
min	8.173617e-08
25%	3.929737e-01
50%	5.659614e-01
75%	6.634218e-01
max	8.549997e-01





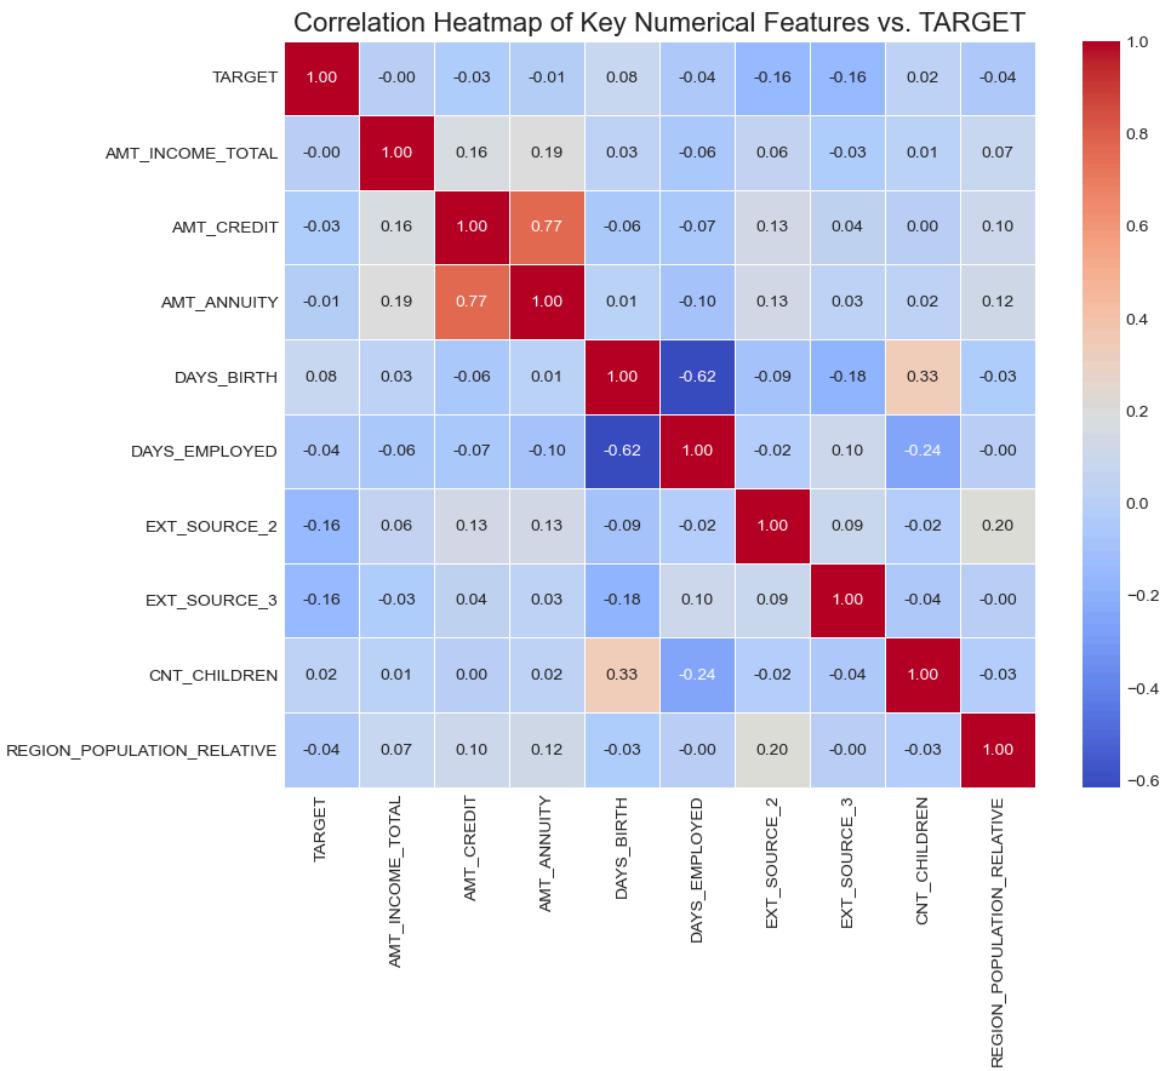
```
In [13]: # Assuming our clean DataFrame is named 'df'
numeric_corr_cols = ['TARGET', 'AMT_INCOME_TOTAL', 'AMT_CREDIT', 'AMT_ANNUITY',
                     'DAYS_BIRTH', 'DAYS_EMPLOYED', 'EXT_SOURCE_2', 'EXT_SOURCE_1',
                     'CNT_CHILDREN', 'REGION_POPULATION_RELATIVE']

correlation_matrix = applicationdf[numeric_corr_cols].corr()

# Printing the table
print("\n--- Correlation Matrix (for identifying top predictors) ---")
print(correlation_matrix.to_string())

# Visualizing the correlation matrix as a heatmap
import matplotlib.pyplot as plt
import seaborn as sns
plt.figure(figsize=(10, 8))
sns.heatmap(correlation_matrix,
            annot=True,
            cmap='coolwarm',
            fmt=".2f",
            linewidths=.5)
plt.title('Correlation Heatmap of Key Numerical Features vs. TARGET', fontsize=14)
plt.show()
```

--- Correlation Matrix (for identifying top predictors) ---						
		TARGET	AMT_INCOME_TOTAL	AMT_CREDIT	AMT_ANNUITY	
DAYS_BIRTH	DAYS_EMPLOYED	EXT_SOURCE_2	EXT_SOURCE_3	CNT_CHILDREN	REGION_POPULATION_RELATIVE	
TARGET		1.000000	-0.003982	-0.030369	-0.012815	
0.078239	-0.044932	-0.160295	-0.155892	0.019187		
-0.037227						
AMT_INCOME_TOTAL		-0.003982	1.000000	0.156870	0.191657	
0.027261	-0.064223	0.060855	-0.030737	0.012882		
0.074796						
AMT_CREDIT		-0.030369	0.156870	1.000000	0.770127	
-0.055436	-0.066838	0.130930	0.036640	0.002145		
0.099738						
AMT_ANNUITY		-0.012815	0.191657	0.770127	1.000000	
0.009443	-0.104329	0.125509	0.026738	0.021377		
0.118418						
DAYS_BIRTH		0.078239	0.027261	-0.055436	0.009443	
1.000000	-0.615864	-0.091947	-0.178527	0.330938		
-0.029582						
DAYS_EMPLOYED		-0.044932	-0.064223	-0.066838	-0.104329	
-0.615864	1.000000	-0.020652	0.101525	-0.239818		
-0.003980						
EXT_SOURCE_2		-0.160295	0.060855	0.130930	0.125509	
-0.091947	-0.020652	1.000000	0.094147	-0.017990		
0.198725						
EXT_SOURCE_3		-0.155892	-0.030737	0.036640	0.026738	
-0.178527	0.101525	0.094147	1.000000	-0.039543		
-0.004507						
CNT_CHILDREN		0.019187	0.012882	0.002145	0.021377	
0.330938	-0.239818	-0.017990	-0.039543	1.000000		
-0.025573						
REGION_POPULATION_RELATIVE	-0.037227		0.074796	0.099738	0.118418	
-0.029582	-0.003980	0.198725	-0.004507	-0.025573		
1.000000						



```
In [14]: from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA # Using PCA to verify clustering in 2D space

# 1. Selecting key numerical features for clustering (excluding TARGET)
clustering_features = ['AMT_INCOME_TOTAL', 'AMT_CREDIT', 'EXT_SOURCE_2', 'EXT_SOURCE_3']
X_cluster = applicationdf[clustering_features]

# 2. Scaling the data (StandardScaler ensures mean=0 and std=1)
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X_cluster)
print("Data scaled successfully.")

# 3. Using PCA to reduce dimensions for visualization purposes (2D plot is better)
pca = PCA(n_components=2)
X_pca = pca.fit_transform(X_scaled)
print(f"PCA reduced data shape: {X_pca.shape}")
```

Data scaled successfully.
 PCA reduced data shape: (307511, 2)

```
In [15]: # Unsupervised Analysis
from sklearn.cluster import KMeans

# 1. Defining the clustering model (K=3)
kmeans = KMeans(n_clusters=3, random_state=42, n_init=10)

# 2. Fitting the model to the scaled data
```

```
kmeans.fit(X_scaled)

# 3. Adding the cluster labels back to the main DataFrame
applicationdf['Cluster'] = kmeans.labels_

print("K-Means clustering completed with 3 clusters.")

# 4. Analyzing the size of each cluster
print("\nCluster Sizes:")
print(applicationdf['Cluster'].value_counts().sort_index())
```

K-Means clustering completed with 3 clusters.

Cluster Sizes:

Cluster	count
0	167683
1	1
2	139827

Name: count, dtype: int64

```
In [16]: # 1. Analyzing the mean features for each cluster to interpret the groups
cluster_profiles = applicationdf.groupby('Cluster')[clustering_features + ['TARGET']]
print("\n--- Cluster Profiles (Mean Feature Values) ---")
print(cluster_profiles)

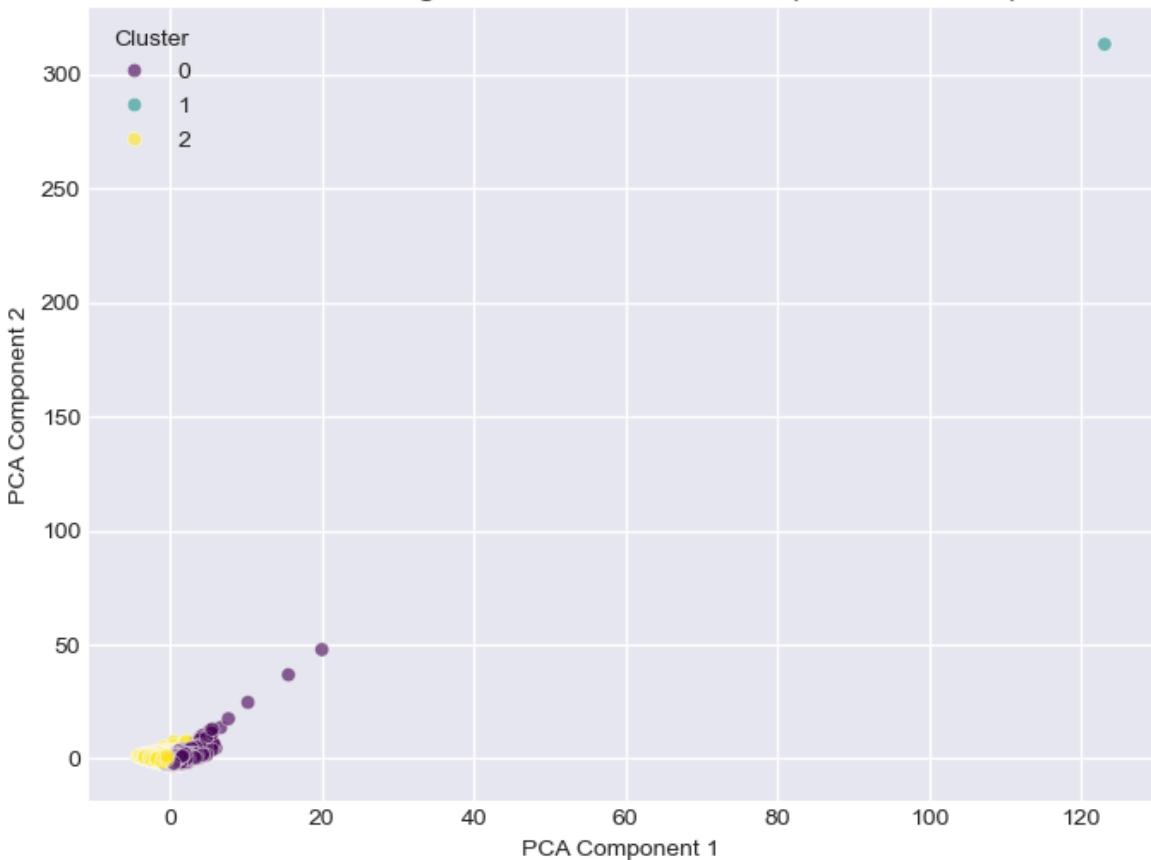
# 2. Visualizing the Clusters in 2D Space (using PCA)
plt.figure(figsize=(8, 6))
sns.scatterplot(x=X_pca[:, 0], y=X_pca[:, 1], hue=applicationdf['Cluster'], palette='viridis')
plt.title('Customer Segmentation via K-Means (PCA Reduced)', fontsize=14)
plt.xlabel('PCA Component 1')
plt.ylabel('PCA Component 2')
plt.show()
```

--- Cluster Profiles (Mean Feature Values) ---

Cluster	AMT_INCOME_TOTAL	AMT_CREDIT	EXT_SOURCE_2	EXT_SOURCE_3	\
0	1.760407e+05	710459.315941	0.582783	0.591805	
1	1.170000e+08	562491.000000	0.113161	0.145543	
2	1.592767e+05	465393.466423	0.432624	0.424425	

Cluster	DAYS_BIRTH	TARGET
0	-18496.962143	0.044125
1	-12615.000000	1.000000
2	-13086.983716	0.124618

Customer Segmentation via K-Means (PCA Reduced)



```
In [17]: # Supervised Analysis

from sklearn.model_selection import train_test_split

# Dropping the SK_ID_CURR (identifier) and the temporary 'Cluster' column before
X = applicationdf.drop(['SK_ID_CURR', 'TARGET', 'Cluster'], axis=1)
Y = applicationdf['TARGET']

# Splitting the data into training (70%) and testing (30%) sets
X_train, X_test, Y_train, Y_test = train_test_split(
    X, Y, test_size=0.3, random_state=42
)

print(f"Data split complete: Train shape = {X_train.shape}, Test shape = {X_test}
```

Data split complete: Train shape = (215257, 182), Test shape = (92254, 182)

```
In [18]: from sklearn.linear_model import LogisticRegression
from sklearn import metrics
import warnings

# Suppressing convergence warnings for cleaner output
warnings.filterwarnings('ignore')

# 1. Defining the model
model = LogisticRegression(solver='liblinear', random_state=42)

# 2. Training the model
model.fit(X_train, Y_train)

# 3. Predicting on the test set
Y_predicted = model.predict(X_test)
```

```
# 4. Evaluating the model (Metrics for the report)
print("\n--- Logistic Regression Model Performance ---")
print("Classification Report:")
print(metrics.classification_report(Y_test, Y_predicted))

print("\nConfusion Matrix:")
print(metrics.confusion_matrix(Y_test, Y_predicted))

# 5. Calculating Accuracy
accuracy = metrics.accuracy_score(Y_test, Y_predicted)
print(f"\nOverall Accuracy: {accuracy:.4f}")
```

--- Logistic Regression Model Performance ---

Classification Report:

	precision	recall	f1-score	support
0	0.92	1.00	0.96	84841
1	0.00	0.00	0.00	7413
accuracy			0.92	92254
macro avg	0.46	0.50	0.48	92254
weighted avg	0.85	0.92	0.88	92254

Confusion Matrix:

```
[[84840    1]
 [ 7413    0]]
```

Overall Accuracy: 0.9196

In []:

[redacted]

In []:

[redacted]